

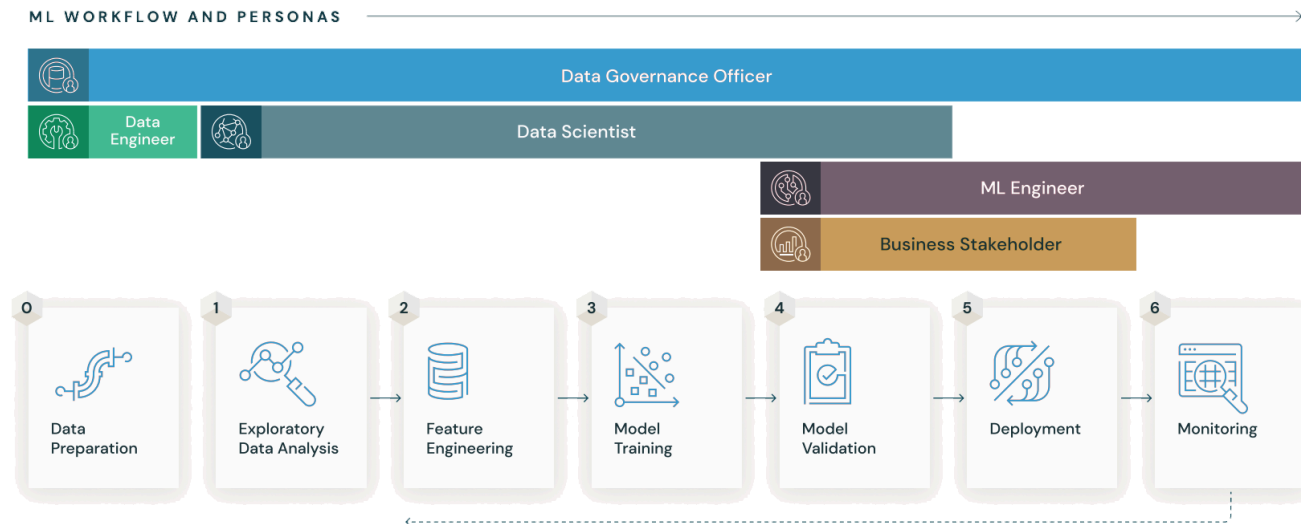
# Introduction to MLflow

**Nima Nooshiri**

16 May 2023

# What is MLflow?

- MLflow is an open source platform for **managing ML and DL projects and experiments**.
- It is specifically built to **optimise the entire model lifecycle**, i.e. from training to production.



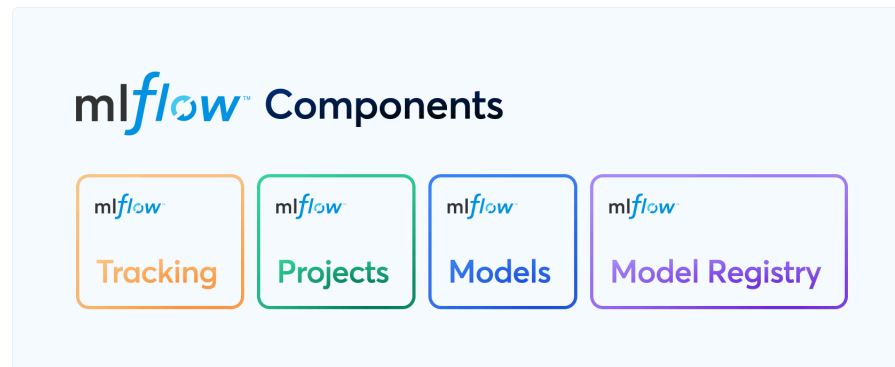
## Comparison of MLflow features with some other experiment tracking and management tools ([resource](#)):

	Neptune	Weights & Biases	Comet	Sacred & Omniboard	MLflow	TensorBoard
<b>Focus</b>	Metadata Storage, Experiment Tracking, Model Registry	Experiment Management	Experiment Management	Experiment Management	Entire Lifecycle	Experiment Management
<b>Price</b>	<ul style="list-style-type: none"> <li>Individual: Free (+ usage above free quota)</li> <li>Academia: Free</li> <li>Team: Paid</li> </ul>	<ul style="list-style-type: none"> <li>Individual: Free (+ usage above free quota)</li> <li>Academia: Free</li> <li>Team: Paid</li> </ul>	<ul style="list-style-type: none"> <li>Individual: Free (+ usage above free quota)</li> <li>Academia: Free</li> <li>Team: Paid</li> </ul>	Free	Free	Free
<b>Standalone component or a part of a broader ML platform?</b>	Standalone component. ML metadata store that focuses on experiment tracking and model registry	Standalone component	Stand-alone tool with community, self-serve and managed deployment options	Omniboard is a web dashboard for the Sacred machine learning experiment management tool	Open-source platform which offers four separate components for experiment tracking, code packaging, model deployment, and model registry	Open source tool which is a part of the TensorFlow ecosystem
<b>Commercial software, open-source software, or a managed cloud service?</b>	Managed cloud service	Managed cloud service	Managed cloud service	Open-source	The standalone product is open-source, while the Databricks managed version is commercial	TensorBoards is open-source, while TensorBoard.dev is available as a free managed cloud service
<b>Hosted version or deployed on-premise?</b>	Tracking is hosted on a managed server, and can also be deployed on-premises and in a public/private cloud server	Yes	Yes, you can deploy Comet on any cloud environment or on-premise	Can be deployed both on-premises and/or on the cloud, but has to be self-managed	Tracking is hosted on a local/remote server (on-prem or cloud). Is also available on a managed server as part of the Databricks platform	TensorBoard is hosted locally. TensorBoard.dev is available on a managed server as a free service
<b>How much do you have to change in your training process?</b>	Minimal. Just a few lines of code needed for tracking	Minimal. Just a few lines of code needed for tracking	Minimal. Few lines of code needed for tracking	Minimal. Only a few lines of code need to be added	Minimal. Few lines of code needed for tracking	Minimal if already using the TensorFlow framework, else significant
<b>Web UI or console-based?</b>	Web UI	Web UI	Web UI	Web UI	Web UI	Web UI

# MLflow Components

MLflow has four primary components:

- **Tracking:** an API and UI for logging parameters, metrics, results and so on when running ML codes.
- **Models:** to manage models from a variety of ML libraries to a variety of model serving and inference platforms.
- **Projects:** to package models in reusable and reproducible form to transfer to production.
- **Model Registry:** a centralized model store managing the full lifecycle of a trained model (model versioning, stage transitions etc).



# Concepts: MLflow Experiment and Run

## Run

- An MLflow run corresponds to a single execution of model code.
- It is a specific ML model with a certain set of hyperparameters, corresponding metrics and results.
- Each run has a unique **Run ID**.

## Experiment

- An MLflow experiment is the primary unit of organisation and access control for MLflow runs.
- An experiment groups together runs for a specific task (all MLflow runs belong to an experiment).
- For each experiment, the results of different runs can be analysed and compared.

## Example

- Create and set an experiment as the active experiment:

```
import mlflow
from mlflow.exceptions import MlflowException

EXPERIMENT_NAME = "IRIS Classifier Exp."

# The experiment must either be specified by name or by ID.
# The experiment name and ID cannot both be specified.

try:
    experiment_id = mlflow.create_experiment(name=EXPERIMENT_NAME)
    experiment = mlflow.get_experiment(experiment_id)
except MlflowException:
    print(f"Experiment {EXPERIMENT_NAME} already exists")
    experiment = mlflow.get_experiment_by_name(EXPERIMENT_NAME)

mlflow.set_experiment(experiment_name=experiment.name)
```

- Run an experiment to train a model:

```
with mlflow.start_run() as run:
    print(f"Starting run with 'run_id': {run.info.run_id}")

    # Do the actual model training
    # ...
```

# MLflow Stores

## Where runs are recorded?

MLflow runs can be recorded:

- to local files,
- to a database,
- remotely to a tracking server.

## How runs and artifacts are recorded?

MLflow uses two components for storage:

1. **Backend store**: persists MLflow entities (runs, parameters, metrics etc).
  - Supported types: *file store* and *database-backed store*.
2. **Artifact store**: persists artifacts (files, models, images, model summary etc).
  - Supported types: *local file paths* and *storage systems* (e.g., AWS S3, GCS)

# MLflow Storage Configuration

## Common scenarios


1. MLflow on **localhost**:
  - Both the backend and artifact stores share a directory on the local filesystem.
2. MLflow on **localhost with SQLite**:
  - Artifacts are stored in a directory on the local filesystem, and MLflow entities are inserted in a SQLite database.
3. MLflow on **localhost with Tracking Server**:
  - Similar to scenario 1 but a tracking server is launched, listening for REST request calls.
4. MLflow with **remote Tracking Server**, backend and artifact stores:
  - The tracking server, backend store, and artifact store reside on remote hosts.



### Example

- Use a SQLite database file as backend store,
- Use a local directory as artifact store,
- Set a tracking server to log runs remotely.

```
$ mlflow server \
  --backend-store-uri="sqlite:///my_database.db" \
  --default-artifact-root="./my_mlruns/" \
  --host="0.0.0.0" \
  --port="5000"
```



# Log Data to Runs

## Basic things to track

- **Parameters:** key-value input parameters using `log_param()` or `log_params()` .
- **Metrics:** key-value metrics using `log_metric()` or `log_metrics()`
- **Artifacts:** local files or directories as artifacts using `log_artifact()` or `log_artifacts()` .

## Example

```
with mlflow.start_run() as run:
    print(f"Starting run with 'run_id': {run.info.run_id}")

    # Log the hyperparameters
    mlflow.log_params(train_params: dict[str: Any])

    # Do the actual model training
    # ...

    # Log the train and test accuracy
    mlflow.log_metric('train_acc', train_acc)
    mlflow.log_metric('test_acc', test_acc)
```

# MLflow Model API

MLflow includes integrations with several common libraries.

## Built-In Model Flavors

- PyTorch
- Keras
- TensorFlow
- Scikit-learn
- XGBoost
- ...

## Key functions

- `mlflow.<model-flavor>.save_model()` : to save the model to a local directory.
- `mlflow.<model-flavor>.log_model()` : to log the model as an artifact in the current run using MLflow Tracking.
- `mlflow.<model-flavor>.load_model()` : to load a model from a local directory or from an artifact in a previous run.

## Example

```
with mlflow.start_run() as run:
    print(f"Starting run with 'run_id': {run.info.run_id}")

    # Log the hyperparameters
    # ...

    # Do the actual model training
    # ...

    # Log the train and test accuracy
    # ...

    # Log the trained model (called "logged model")
    mlflow.pytorch.log_model(pytorch_model=my_model, artifact_path='pytorch_models')
```

# MLflow Model Registry Concept

## Model

- An MLflow Model is created from an experiment or run that is logged with one of the model flavor's methods:

```
mlflow.<model-flavor>.log_model()
```

- Once logged, this model can then be registered with the Model Registry.

## Registered Model

- An MLflow Model can be registered with the Model Registry.
- A registered model has a unique name, contains versions, associated transitional stages, model lineage, and other metadata.

# Register a Model

There are three ways to add a model to the registry:

1. Using the `mlflow.<model-flavor>.log_model(..., registered_model_name=<name>)` method.
2. Using the `mlflow.register_model()` method.
3. Using the `create_registered_model()` method.

## Example

```
import mlflow
from mlflow.tracking import MlflowClient
from mlflow.exceptions import RestException

MODEL_URI = "runs:/3f2a871bfe944121a2372586ec4ae17f/pytorch_models"
REGISTERED_MODEL_NAME = "IRIS_CLF_MLP"

# Create registered model (registry point)
client = MlflowClient()

try:
    registered_model = client.create_registered_model(REGISTERED_MODEL_NAME)
except RestException:
    print(f"Model already exists in the registry: '{REGISTERED_MODEL_NAME}'")

# Register experiment run to that model
model_version = mlflow.register_model(MODEL_URI, REGISTERED_MODEL_NAME)
```

## Load a Model

- To load a previously logged model for inference or development, use `mlflow.<model-flavour>.load_model(model_uri)`.
- The location of the model, `model_uri`, can be one of the following:
  - A run-relative path: `runs:/<run-id>/<run-relative-path-to-model>`,
  - A registered-model path: `models:/<model-name>/<model-version>` or `models:/<model-name>/<stage>`.

## Example: Classification with PyTorch and MLflow

This example demonstrates training a classification model on the IRIS dataset, creating the model with PyTorch, logging the model to MLflow, registering the trained model, and loading it back for inference.

### Step 1: Set up the tracking server

**NOTE:** In order to use model registry functionality, we must run the server using a database-backed backend store.

```
$ mlflow server \
  --backend-store-uri="sqlite:///mlflow.db" \
  --default-artifact-root="./mlruns/" \
  --host="0.0.0.0" \
  --port="5000"
```

Now, we are able to see the tracking URI in the browser at <http://0.0.0.0:5000>.



## Step 2: Train and Log the Model

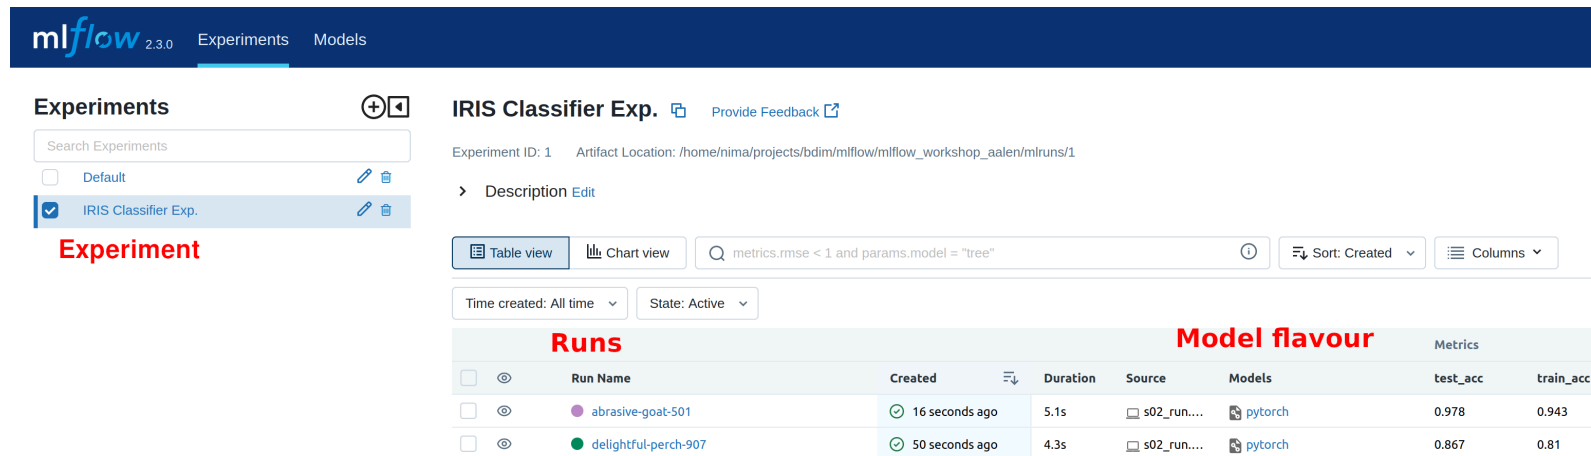
To start a run to train a model and then log the trained model:

```
$ python s02_train_and_log.py --lr=0.005 --dropout-prob=0.1 --num-epochs=150
```

The tracking URI and the experiment name are already set (hard-coded) in the file.

To use a different tracking URI and/or change the experiment name, they can be changed in the python file:

```
TRACKING_SERVER_URI = "http://0.0.0.0:5000"
EXPERIMENT_NAME = "Classification_Experiment"
```



The screenshot shows the mlflow 2.3.0 Experiments page. The left sidebar lists experiments, with 'IRIS Classifier Exp.' selected. The main area shows details for this experiment, including its ID (1) and artifact location. Below this, there are tabs for 'Table view' and 'Chart view', a search bar, and filters for 'Time created' and 'State'. A table of runs is displayed, showing two runs: 'abrasive-goat-501' and 'delightful-perch-907'. The table includes columns for Run Name, Created, Duration, Source, Models, and Metrics (test\_acc, train\_acc).

Runs		Model flavour			Metrics	
Run Name	Created	Duration	Source	Models	test_acc	train_acc
abrasive-goat-501	16 seconds ago	5.1s	s02_run....	pytorch	0.978	0.943
delightful-perch-907	50 seconds ago	4.3s	s02_run....	pytorch	0.867	0.81

## Step 3: Register the Logged Model

```
$ python s03_register \
  --logged-model-uri='runs:/0bc56a73394c4a65bdd3b9d69e3aa0f8/pytorch_models' \
  --registered-model-name='MLP_Classifier'
```

mlflow 2.3.0 Experiments Models GitHub Docs

IRIS Classifier Exp. >  
**abrasive-goat-501**

Run ID: 0bc56a73394c4a65bdd3b9d69e3aa0f8 Date: 2023-05-16 07:42:05 Source: s02\_run.py User: nima Duration: 5.1s

Status: FINISHED **Run ID (unique)** Lifecycle Stage: active

> Description Edit

> Parameters (3)

> Metrics (2)

> Tags (1)

▼ Artifacts

pytorch\_models

data

MLmodel

conda.yaml

python\_env.yaml

requirements.txt

s02\_run.py

Full Path:/home/nima/projects/bdim/mlflow/mlflow\_workshop\_aalen/miruns/1/0bc56a73394c4a65bdd3b9d69e3aa0f8/artifacts/pytorch\_models

Register Model

### MLflow Model

The code snippets below demonstrate how to make predictions using the logged model. You can also [register it to the model registry](#) to version control

#### Model schema

Input and output schema for your model. [Learn more](#)

Name	Type
Inputs (1)	
Outputs (1)	

#### Make Predictions

Predict on a Spark DataFrame:

```
import mlflow
from pyspark.sql.functions import struct, col
logged_model = 'runs:/0bc56a73394c4a65bdd3b9d69e3aa0f8/pytorch_models'

# Load model as a Spark UDF. Override result_type if the model does not return double values.
loaded_model = mlflow.pyfunc.spark_udf(spark, model_uri=logged_model, result_type='double')

# Predict on a Spark DataFrame.
df.withColumn('predictions', loaded_model(struct(*map(col, df.columns))))
```

**The URI of the logged model**

mlflow2.3.0

Experiments

Models

Experiments

Search Experiments

Default

IRIS Classifier Exp.

IRIS Classifier Exp.

Provide Feedback

Experiment ID: 1

Artifact Location: /home/nima/projects/bdim/mlflow/mlflow\_workshop\_aalen/mlruns/1

Description

Edit

Table view

Chart view

metrics.rmse < 1 and params.model = "tree"

Sort: Created

Columns

Time created: All time

State: Active

Registered model with a new icon and name

							Metrics	
	Run Name	Created	Duration	Source	Models		test_acc	train_acc
	abrasive-goat-501	4 minutes ago	5.1s	s02_run....	IRIS_CLF_M.../1		0.978	0.943
	delightful-perch-907	5 minutes ago	4.3s	s02_run....	pytorch		0.867	0.81

mlflow2.3.0

Experiments

Models

Registered Models

Share and manage machine learning models. [Learn more](#)

Create Model

Name	Latest Version	Staging	Production
IRIS_CLF_MLP	Version 1	—	—

## Step 4: Use the Model for Inference

Now, we can select a model, load it and make a prediction.

```
# Use model version 1
$ python s04_predict --logged-model-uri="models:/MLP_Classifier/1"
```

We can also specify the model **tag** instead of model version. For example, if we want to use a model with tag `champion`, we need to replace `/1` with `@champion` in the model URI above.

## Resources

- [MLflow documentation](#).
- [MLflow examples](#).
- [Neptune.ai blog](#)